

# Class 7: Machine Learning 1

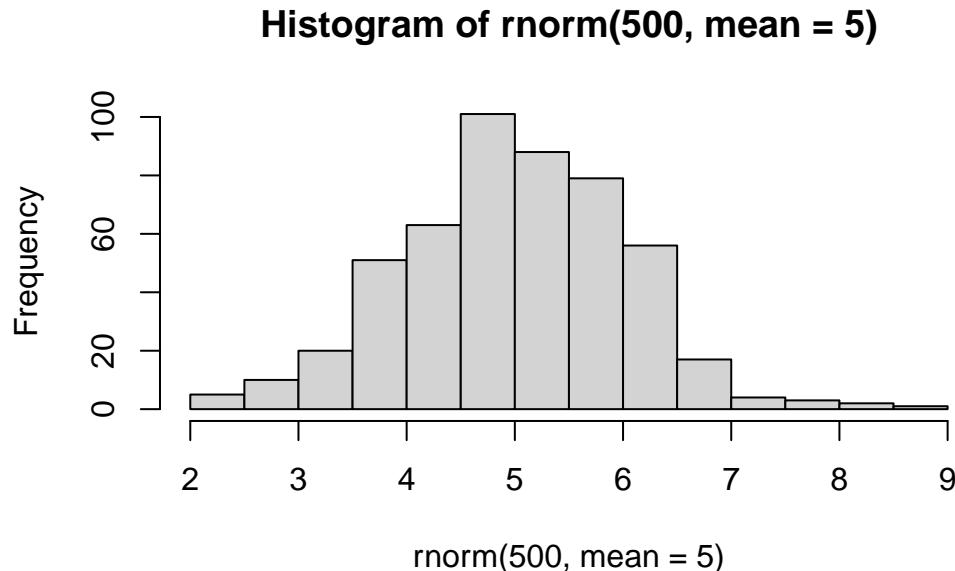
Joseph Girgiss (PID: A17388247)

Today we will explore some fundamental machine learning methods including clustering and dimensionality reduction.

## K-means clustering

To see how this works let's first makeup some data to cluster where we know what the answer should be. We can use the `rnorm()` function to help here:

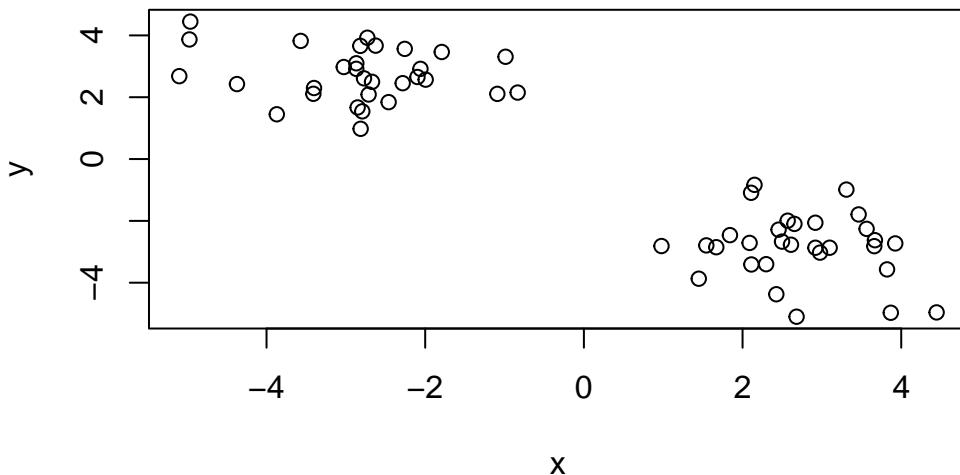
```
hist(rnorm(500, mean = 5))
```



```
x <- c(rnorm(30, mean = -3), rnorm(30, mean =3 ))
y <- rev(x)
```

```
x <- cbind(x,y)
```

plot(x)



The function for K-means clustering in “base” R is `kmeans()`

```
k <- kmeans(x, centers = 2)  
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	2.725280	-2.835451
2	-2.835451	2.725280

Clustering vector:

[39] 1

```
Within cluster sum of squares by cluster:  
[1] 53.60948 53.60948  
(between_SS / total_SS =  89.6 %)
```

## Available components:

```
[1] "cluster"        "centers"        "totss"          "withinss"        "tot.withinss"  
[6] "betweenss"     "size"           "iter"           "ifault"
```

To get at the results of the returned list object we can use the dollar \$ syntax.

Q. How many points are in each cluster?

k\$size

[1] 30 30

Q. What ‘component’ of your result object details

- cluster size?
  - cluster assignment/membership?
  - cluster cent

k\$cluster

k\$center

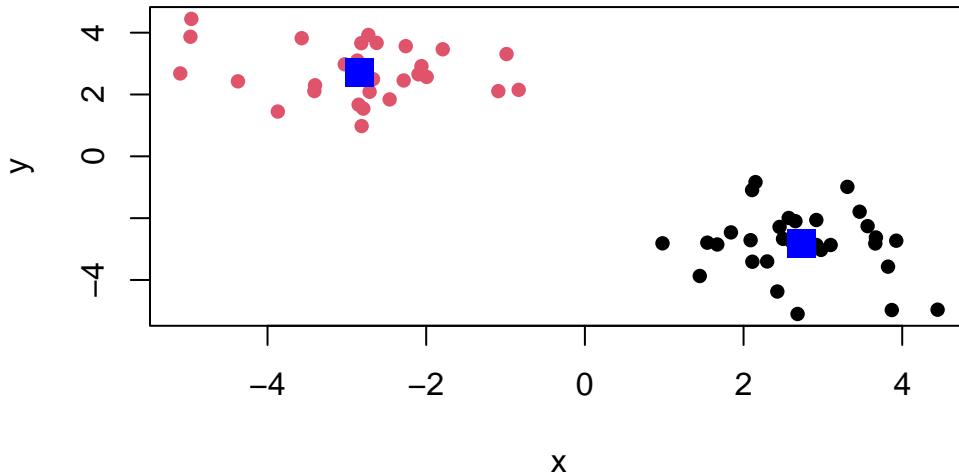
```

          x           y
1  2.725280 -2.835451
2 -2.835451  2.725280

```

Q. Make a clustering results figure of the data colored by cluster membership and show cluster centers.

```
plot(x, col = k$cluster, pch = 16)
points(k$centers, col = "blue", pch = 15, cex = 2)
```

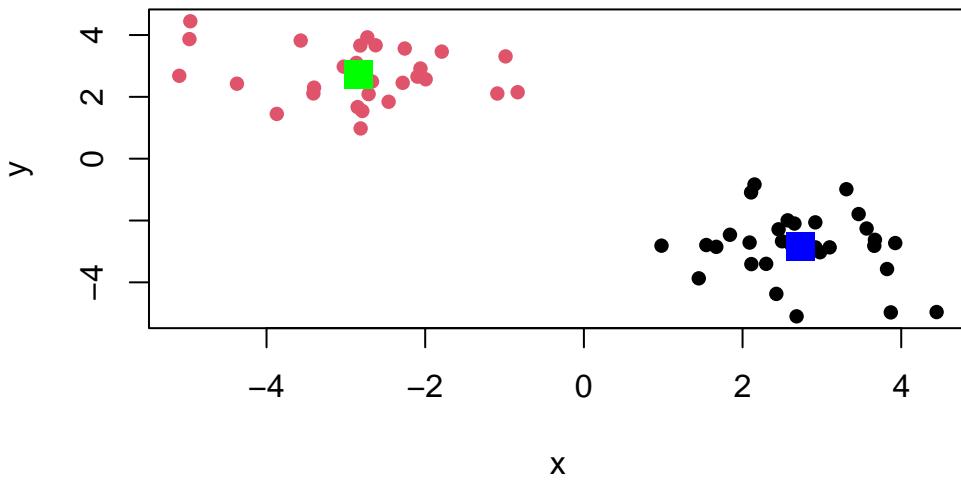


K-means clustering is very popular as it is very fast and relatively straight forward: it takes numeric data as input and returns the cluster membership vector etc.

The “issue” is we tell `kmeans()` how many clusters we want!

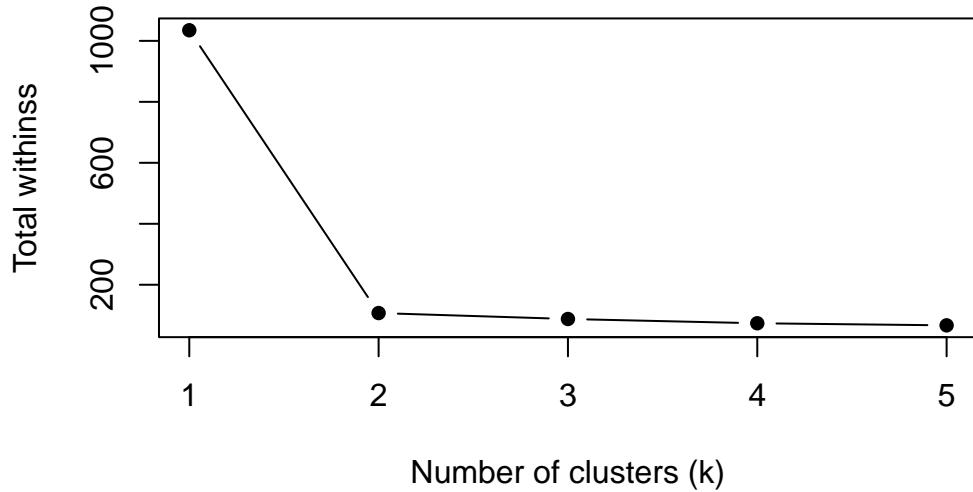
Q. Run k-means again and cluster 4 groups/clusters and plot the results as we did above?

```
k4 <- kmeans(x, centers = 4)
plot(x, col = k$cluster, pch = 16)
points(k$centers, col = c("blue", "green", "purple", "red"), pch = 15, cex = 2)
```

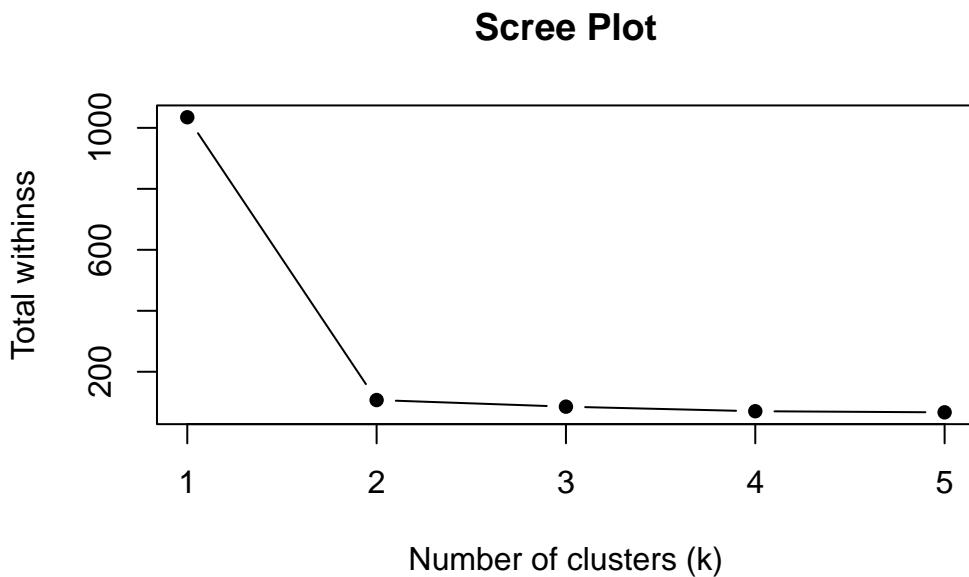


```
wss <- numeric()
for (k in 1:5) {
  wss[k] <- kmeans(x, centers = k)$tot.withinss
}
plot(1:5, wss, type = "b", xlab = "Number of clusters (k)", ylab = "Total withinss", main = "Elbow Method Plot")
```

## Scree Plot



```
n <- NULL
for(i in 1:5) {
  n <- c(n, kmeans(x, centers = i)$tot.withinss)
}
plot(n, typ="b", xlab = "Number of clusters (k)", ylab = "Total withinss", main = "Scree Plot")
```



The elbow point is the location on a scree plot where the total within-cluster sum of squares drops sharply as the number of clusters increases, but then levels off. This point represents the optimal number of clusters.

### Next Hierarchical Clustering

The main “base” R function for Hierarchical Clustering is called `hclust()`. Here we can’t just input our data we need to first calculate a distance matrix (e.g `dist()`) for our data and use this as input to `hclust()`

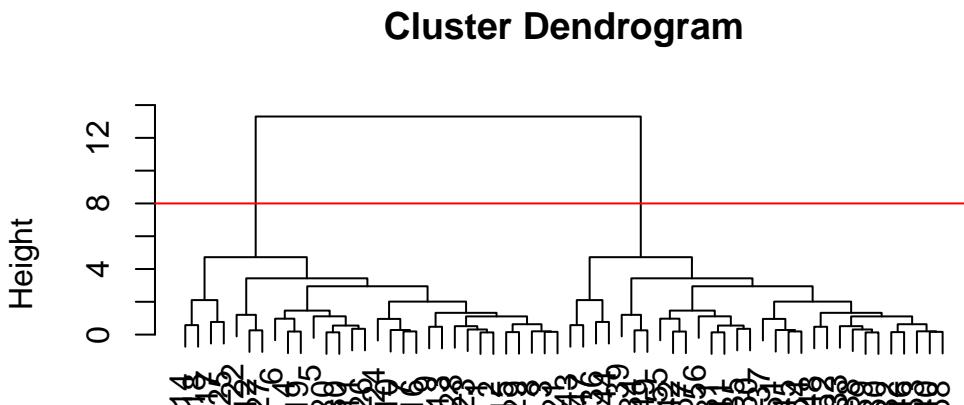
```
d <- dist(x)
hc <- hclust(d)
hc
```

```
Call:
hclust(d = d)

Cluster method : complete
Distance       : euclidean
Number of objects: 60
```

There is a plot method for hclust results. Let's try it!

```
plot(hc)
abline(h=8, col = "red")
```



```
d  
hclust (*, "complete")
```

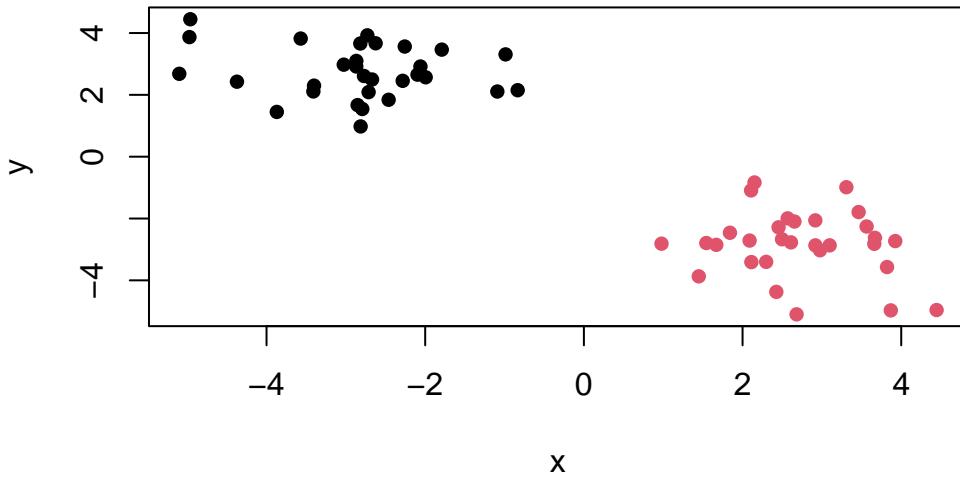
```
cutree(hc, h=8)
```

To get our cluster “membership” vector (i.e. our main clustering result) we can “cut” the tree at a given height that yields a given “ $k$ ” groups.

```
grps <- cutree(hc, k=2)
```

Q. Plot the data with our hclust result coloring

```
plot(x, col = grps, pch = 16)
```



## Principal Component Analysis (PCA)

### PCA of UK food data

Import food data from an online CSV file:

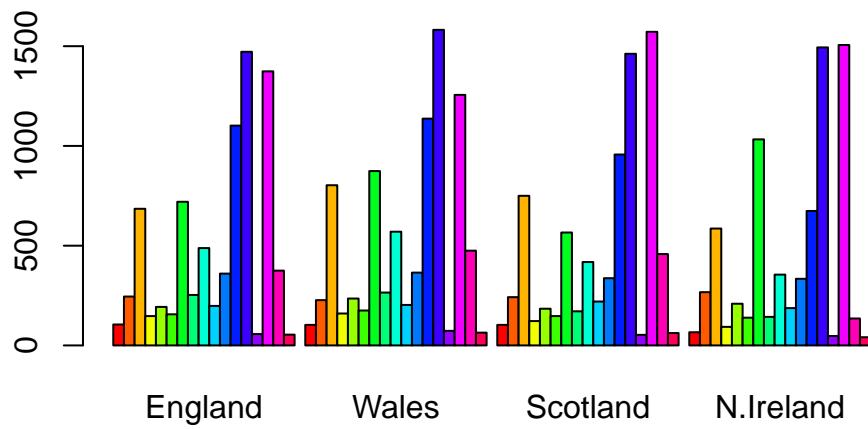
```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names = 1)
x
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187

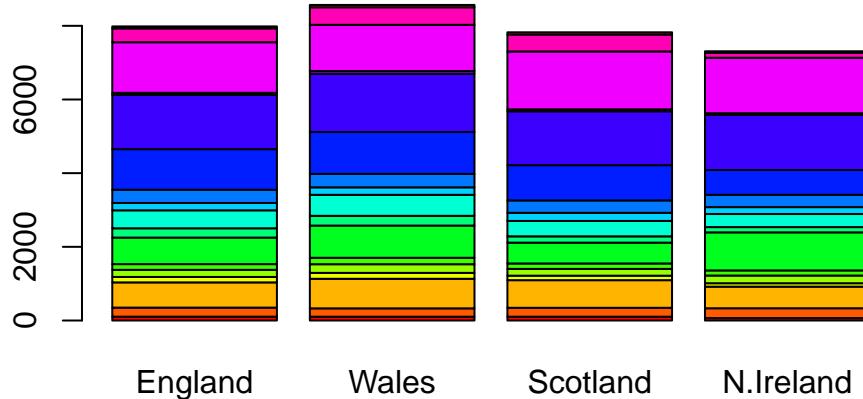
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41

Some base figures:

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```

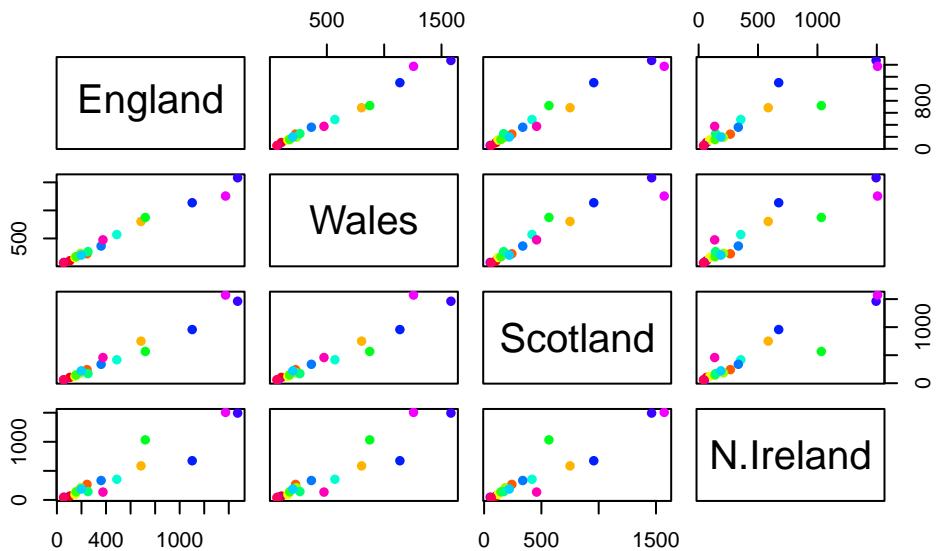


```
barplot(as.matrix(x), col=rainbow(nrow(x)))
```



There is one plot that can be useful for small datasets:

```
pairs(x, col=rainbow(nrow(x)), pch=16)
```



Main point: it can be difficult to spot major trends and patterns even in relatively small multivariate datasets (here we only have 17 dimensions, typically we have 1000s).

## PCA to the rescue

The main function in “base” R for PCA is called `prcomp()`

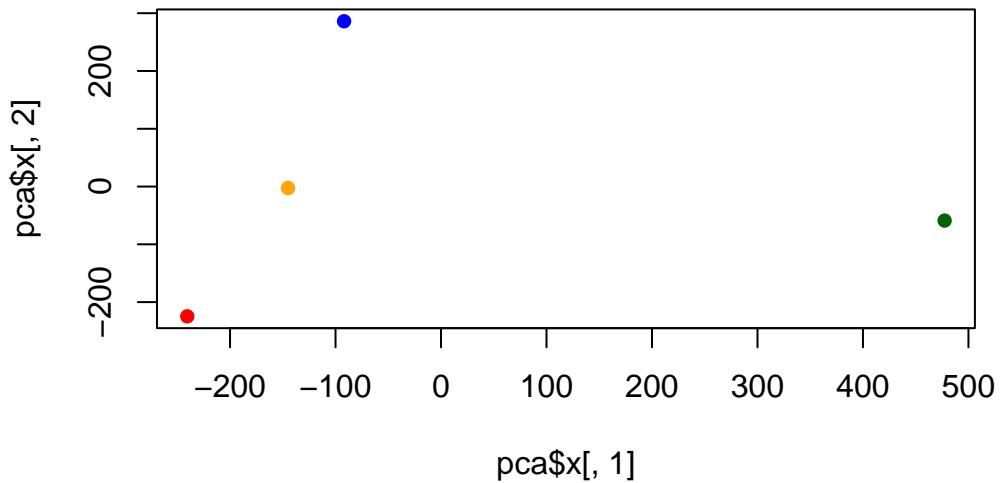
I will take the transpose of our data so the “foods” are in the columns:

```
pca <- prcomp( t(x) )
summary(pca)
```

Importance of components:

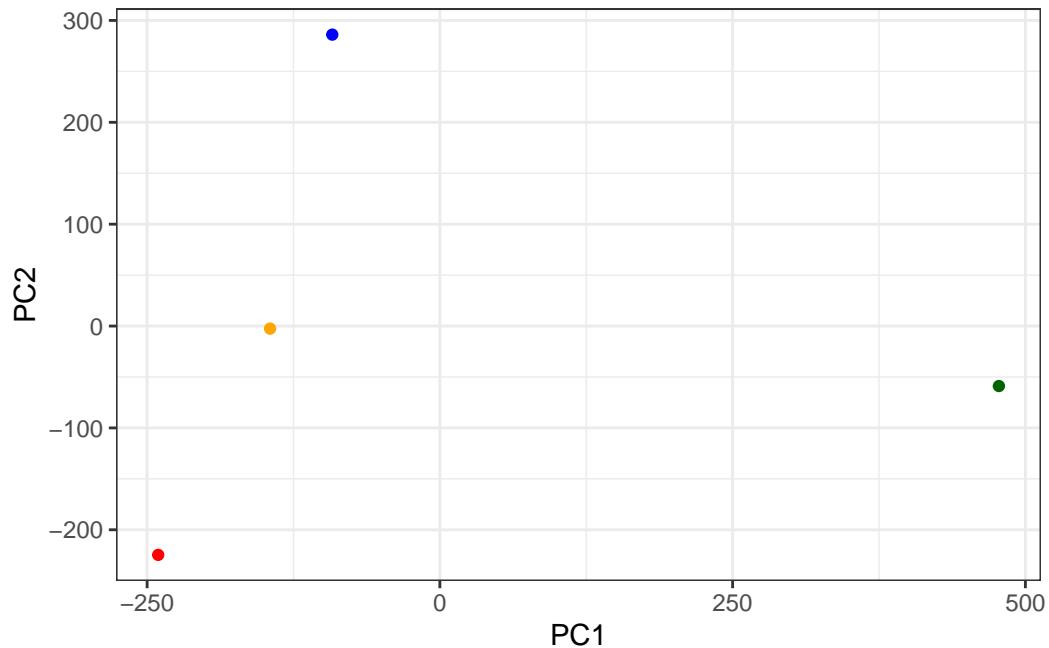
	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

```
cols <- c("orange", "red", "blue", "darkgreen")
plot(pca$x[,1], pca$x[,2], col=cols, pch=16)
```

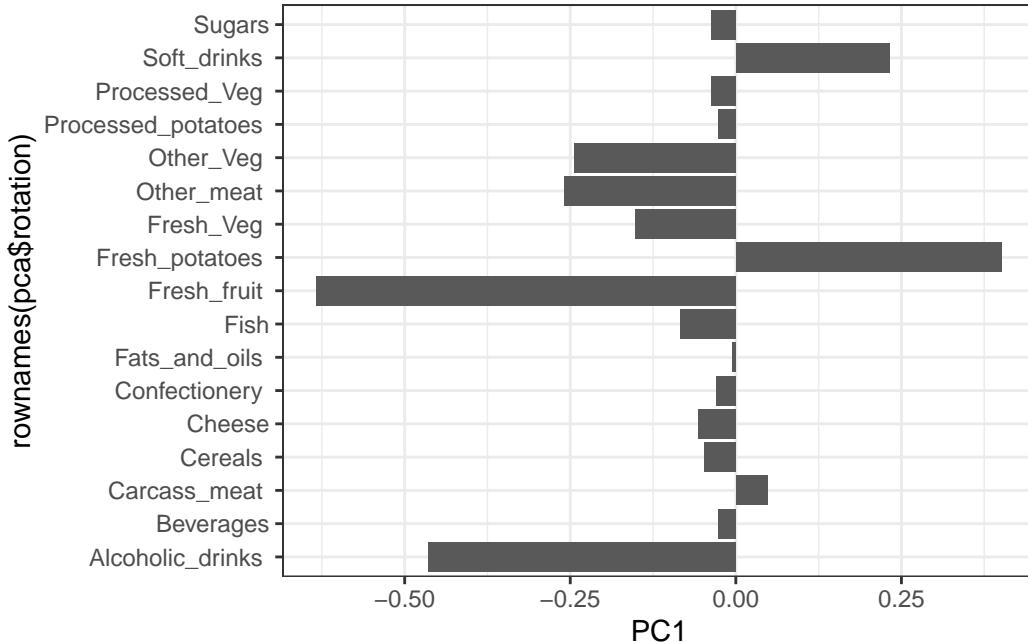


```
library(ggplot2)

ggplot(pca$x) +
  aes(PC1, PC2) +
  geom_point(col=cols) +
  theme_bw()
```



```
ggplot(pca$rotation) +  
  aes(PC1, rownames(pca$rotation)) +  
  geom_col() +  
  theme_bw()
```



PCA looks super useful and we will come back to describe this further next time.