

Class 8: Mini Project

Joseph Girgiss (PID: A17388247)

Table of contents

Background	1
Data input	2
Exploratory data analysis	3
Principal Component Analysis (PCA)	4
Interpreting PCA results	7
Variance explained	9
Communicating PCA results	12
Hierarchical clustering	13
K-means clustering	16
Combining Methods (PCA and Clustering)	16
Prediction	19

Background

The goal of this mini-project is for you to explore a complete analysis using the unsupervised learning techniques covered in class. You'll extend what you've learned by combining PCA as a preprocessing step to clustering using data that consist of measurements of cell nuclei of human breast masses. This expands on our RNA-Seq analysis from last day.

The data itself comes from the Wisconsin Breast Cancer Diagnostic Data Set first reported by K. P. Benne and O. L. Mangasarian: "Robust Linear Programming Discrimination of Two Linearly Inseparable Sets".

Values in this data set describe characteristics of the cell nuclei present in digitized images of a fine needle aspiration (FNA) of a breast mass.

Data input

Data was downloaded from the class website as a CSV file.

```
# Save your input data file into your Project directory  
fna.data <- "WisconsinCancer.csv"  
  
# Complete the following code to input the data and store as wisc.df  
wisc.df <- read.csv(fna.data, row.names=1)  
  
# First 6 rows of data.frame  
head(wisc.df)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
84358402	M	20.29	14.34	135.10	1297.0
843786	M	12.45	15.70	82.57	477.1
	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean	
842302	0.11840	0.27760	0.3001	0.14710	
842517	0.08474	0.07864	0.0869	0.07017	
84300903	0.10960	0.15990	0.1974	0.12790	
84348301	0.14250	0.28390	0.2414	0.10520	
84358402	0.10030	0.13280	0.1980	0.10430	
843786	0.12780	0.17000	0.1578	0.08089	
	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419		0.07871	1.0950	0.9053
842517	0.1812		0.05667	0.5435	0.7339
84300903	0.2069		0.05999	0.7456	0.7869
84348301	0.2597		0.09744	0.4956	1.1560
84358402	0.1809		0.05883	0.7572	0.7813
843786	0.2087		0.07613	0.3345	0.8902
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
842302	153.40	0.006399	0.04904	0.05373	0.01587
842517	74.08	0.005225	0.01308	0.01860	0.01340
84300903	94.03	0.006150	0.04006	0.03832	0.02058
84348301	27.23	0.009110	0.07458	0.05661	0.01867
84358402	94.44	0.011490	0.02461	0.05688	0.01885
843786	27.19	0.007510	0.03345	0.03672	0.01137
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst	

842302	0.03003	0.006193	25.38	17.33
842517	0.01389	0.003532	24.99	23.41
84300903	0.02250	0.004571	23.57	25.53
84348301	0.05963	0.009208	14.91	26.50
84358402	0.01756	0.005115	22.54	16.67
843786	0.02165	0.005082	15.47	23.75
	perimeter_worst	area_worst	smoothness_worst	compactness_worst
842302	184.60	2019.0	0.1622	0.6656
842517	158.80	1956.0	0.1238	0.1866
84300903	152.50	1709.0	0.1444	0.4245
84348301	98.87	567.7	0.2098	0.8663
84358402	152.20	1575.0	0.1374	0.2050
843786	103.40	741.6	0.1791	0.5249
	concavity_worst	concave.points_worst	symmetry_worst	
842302	0.7119	0.2654	0.4601	
842517	0.2416	0.1860	0.2750	
84300903	0.4504	0.2430	0.3613	
84348301	0.6869	0.2575	0.6638	
84358402	0.4000	0.1625	0.2364	
843786	0.5355	0.1741	0.3985	
	fractal_dimension_worst			
842302		0.11890		
842517		0.08902		
84300903		0.08758		
84348301		0.17300		
84358402		0.07678		
843786		0.12440		

```
# We can use -1 here to remove the first column (remove patient ID)
wisc.data <- wisc.df[,-1]
```

We will create a diagnosis vector for use later when we compare information.

```
diagnosis <- factor(wisc.df$diagnosis)
```

Exploratory data analysis

Q1. How many observations are in this dataset?

There are 569 observations/patients in data set.

Q2. How many of the observations have a malignant diagnosis?

```
table(wisc.df$diagnosis)
```

```
B      M  
357  212
```

212 malignant diagnoses, and 357 benign diagnoses.

Q3. How many variables/features in the data are suffixed with _mean?

```
length(grep("_mean$", colnames(wisc.data)))
```

```
[1] 10
```

Principal Component Analysis (PCA)

The `prcomp()` function to do PCA has a `scale=FALSE` default. We want `scale=TRUE` so our analysis is not dominated by columns/variables in our dataset that have high standard deviation and mean when compared to others just because the units of measurement are different.

```
# Check column means and standard deviations  
colMeans(wisc.data)
```

radius_mean	texture_mean	perimeter_mean
1.412729e+01	1.928965e+01	9.196903e+01
area_mean	smoothness_mean	compactness_mean
6.548891e+02	9.636028e-02	1.043410e-01
concavity_mean	concave.points_mean	symmetry_mean
8.879932e-02	4.891915e-02	1.811619e-01
fractal_dimension_mean	radius_se	texture_se
6.279761e-02	4.051721e-01	1.216853e+00
perimeter_se	area_se	smoothness_se
2.866059e+00	4.033708e+01	7.040979e-03
compactness_se	concavity_se	concave.points_se
2.547814e-02	3.189372e-02	1.179614e-02
symmetry_se	fractal_dimension_se	radius_worst
2.054230e-02	3.794904e-03	1.626919e+01
texture_worst	perimeter_worst	area_worst
2.567722e+01	1.072612e+02	8.805831e+02

```

smoothness_worst      compactness_worst      concavity_worst
1.323686e-01          2.542650e-01          2.721885e-01
concave.points_worst   symmetry_worst        fractal_dimension_worst
1.146062e-01          2.900756e-01          8.394582e-02

```

```
apply(wisc.data, 2, sd)
```

radius_mean	texture_mean	perimeter_mean
3.524049e+00	4.301036e+00	2.429898e+01
area_mean	smoothness_mean	compactness_mean
3.519141e+02	1.406413e-02	5.281276e-02
concavity_mean	concave.points_mean	symmetry_mean
7.971981e-02	3.880284e-02	2.741428e-02
fractal_dimension_mean	radius_se	texture_se
7.060363e-03	2.773127e-01	5.516484e-01
perimeter_se	area_se	smoothness_se
2.021855e+00	4.549101e+01	3.002518e-03
compactness_se	concavity_se	concave.points_se
1.790818e-02	3.018606e-02	6.170285e-03
symmetry_se	fractal_dimension_se	radius_worst
8.266372e-03	2.646071e-03	4.833242e+00
texture_worst	perimeter_worst	area_worst
6.146258e+00	3.360254e+01	5.693570e+02
smoothness_worst	compactness_worst	concavity_worst
2.283243e-02	1.573365e-01	2.086243e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
6.573234e-02	6.186747e-02	1.806127e-02

```
# Perform PCA on wisc.data by completing the following code
wisc.pr <- prcomp(wisc.data, scale = T)
```

```
# Look at summary of results
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624

```

Proportion of Variance 0.01589 0.0139 0.01169 0.0098 0.00871 0.00805 0.00523
Cumulative Proportion 0.92598 0.9399 0.95157 0.9614 0.97007 0.97812 0.98335
PC15 PC16 PC17 PC18 PC19 PC20 PC21
Standard deviation 0.30681 0.28260 0.24372 0.22939 0.22244 0.17652 0.1731
Proportion of Variance 0.00314 0.00266 0.00198 0.00175 0.00165 0.00104 0.0010
Cumulative Proportion 0.98649 0.98915 0.99113 0.99288 0.99453 0.99557 0.9966
PC22 PC23 PC24 PC25 PC26 PC27 PC28
Standard deviation 0.16565 0.15602 0.1344 0.12442 0.09043 0.08307 0.03987
Proportion of Variance 0.00091 0.00081 0.0006 0.00052 0.00027 0.00023 0.00005
Cumulative Proportion 0.99749 0.99830 0.9989 0.99942 0.99969 0.99992 0.99997
PC29 PC30
Standard deviation 0.02736 0.01153
Proportion of Variance 0.00002 0.00000
Cumulative Proportion 1.00000 1.00000

```

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)?

Proportion of Variance for PC1 is 0.4427.

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

3 PCs

```
which(summary(wisc.pr)$importance["Cumulative Proportion", ] >= 0.70)[1]
```

```

PC3
3

```

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

7 PCs.

```
which(summary(wisc.pr)$importance["Cumulative Proportion", ] >= 0.90)[1]
```

```

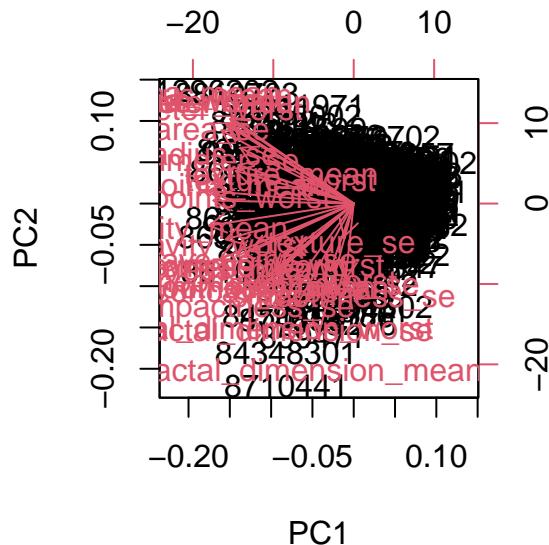
PC7
7

```

Interpreting PCA results

The main PC result figure is called a “score plot” or “PC plot” or “ordination plot”

```
biplot(wisc.pr)
```

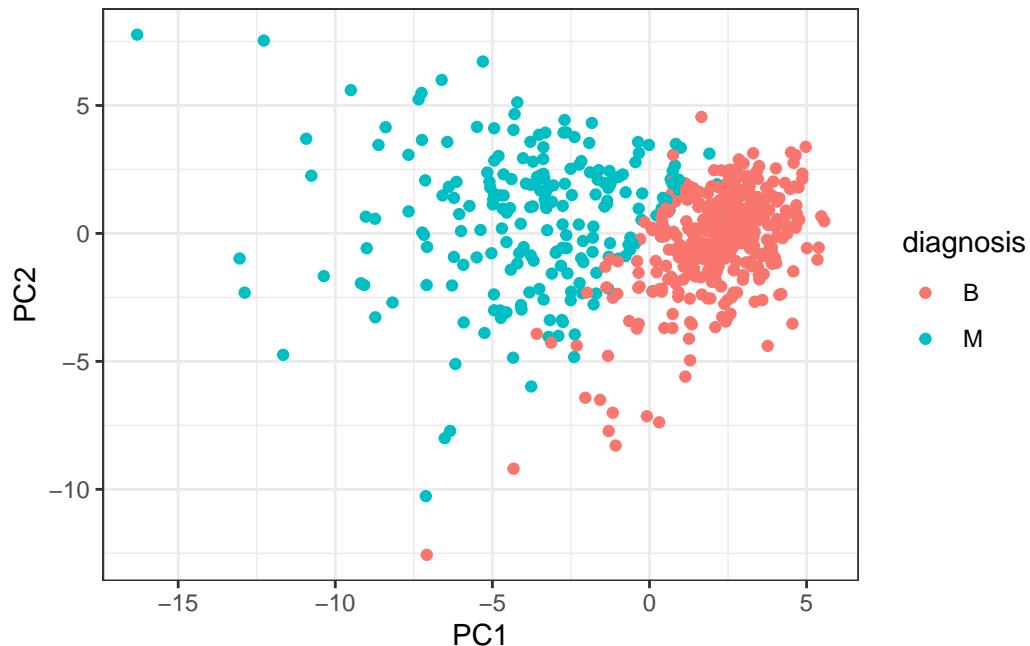


Q7. What stands out to you about this plot? Is it easy or difficult to understand?
Why?

`biplot(wisc.pr)` is very difficult to read/understand. Too many points.

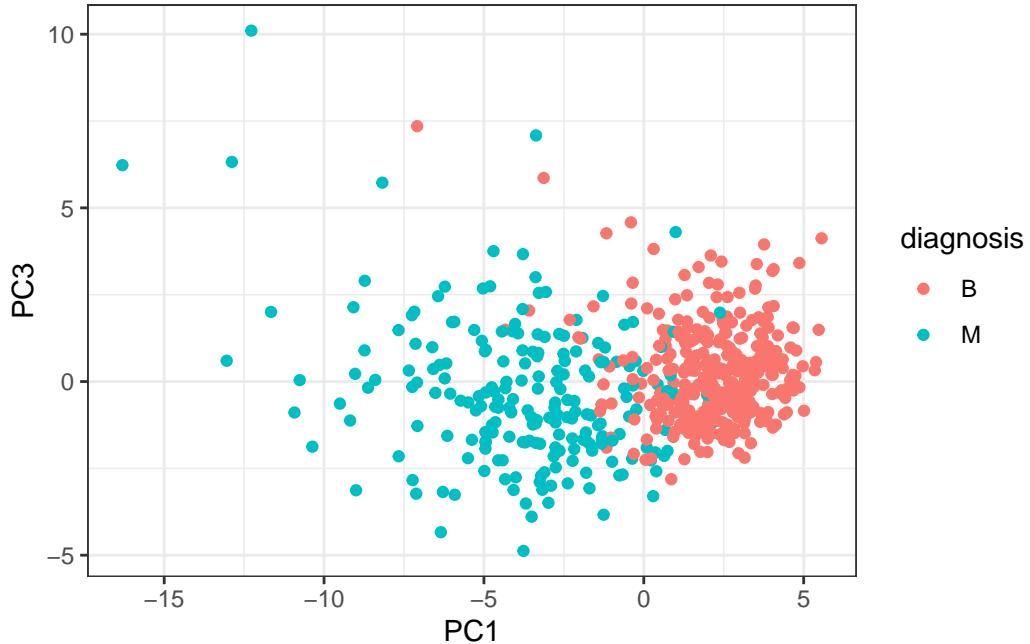
```
library(ggplot2)

#Scatter plot observations by components 1 and 2
ggplot(wisc.pr$x) +
  aes(PC1, PC2, col = diagnosis) +
  geom_point() +
  theme_bw()
```



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

```
ggplot(wisc.pr$x) +
  aes(PC1, PC3, col = diagnosis) +
  geom_point() +
  theme_bw()
```



The values in the PC1 and PC3 plots tend to be more negative than the PC1 and PC2 plot (therefore, PC1 and PC2 plot have more positive values than the PC1 and PC3 plot). There is more overlap between PC1 and PC3 than PC1 and PC2.

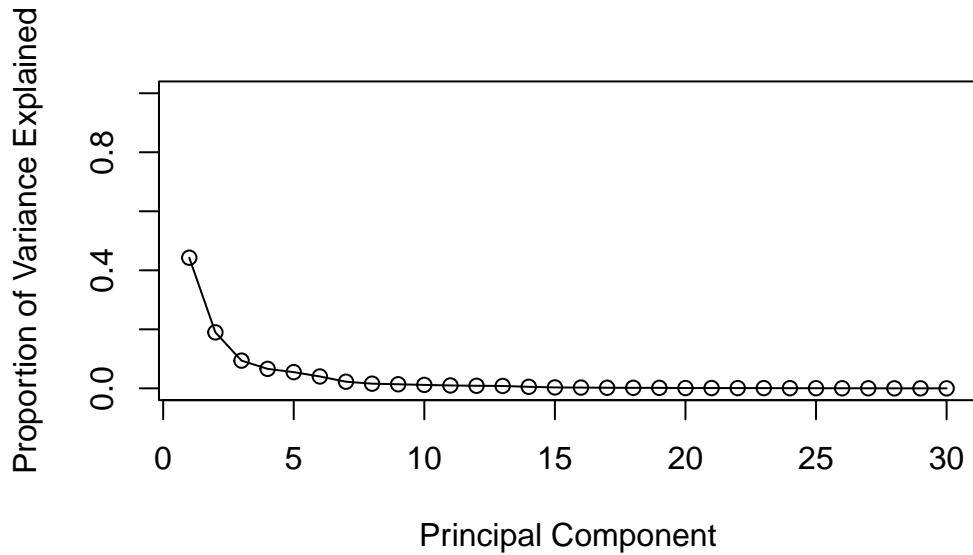
Variance explained

```
# Calculate variance of each component
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

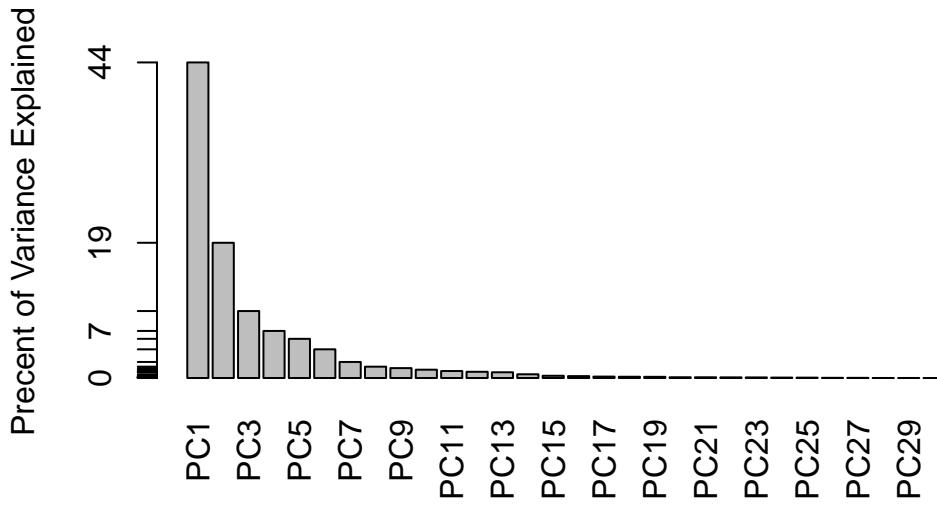
```
[1] 13.281608 5.691355 2.817949 1.980640 1.648731 1.207357
```

```
# Variance explained by each principal component: pve
pve <- pr.var / sum(pr.var)

# Plot variance explained for each principal component
plot(pve, xlab = "Principal Component",
      ylab = "Proportion of Variance Explained",
      ylim = c(0, 1), type = "o")
```



```
# Alternative scree plot of the same data, note data driven y-axis  
barplot(pve, ylab = "Percent of Variance Explained",  
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)  
axis(2, at=pve, labels=round(pve,2)*100 )
```

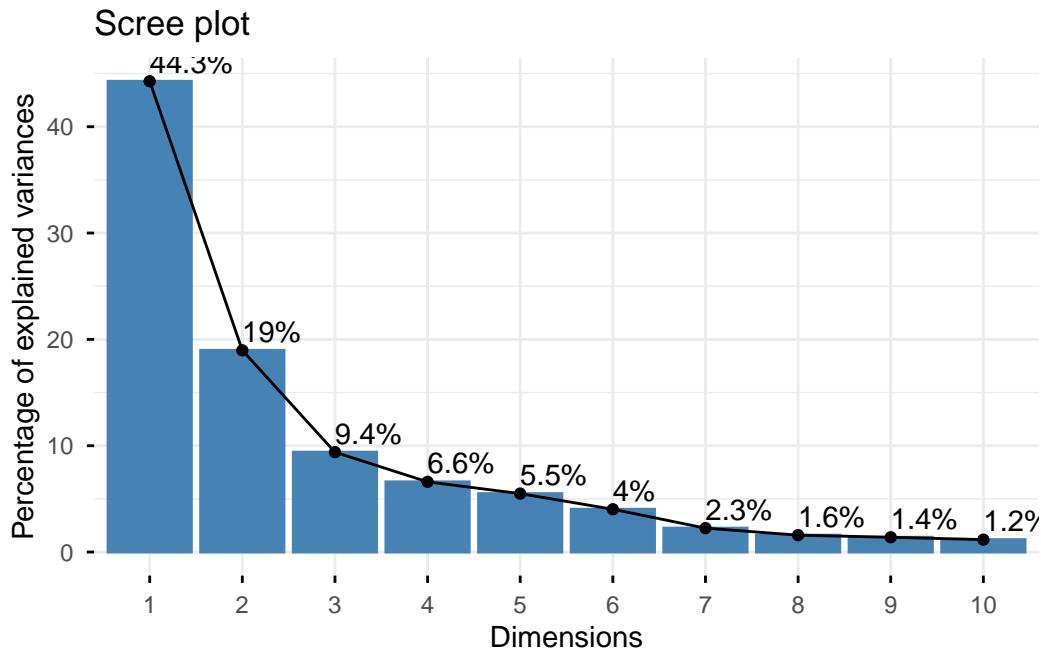


```
## ggplot based graph
#install.packages("factoextra")
library(factoextra)
```

Welcome! Want to learn more? See two factoextra-related books at <https://goo.gl/ve3WBa>

```
fviz_eig(wisc.pr, addlabels = TRUE)
```

Warning in geom_bar(stat = "identity", fill = barfill, color = barcolor, :
Ignoring empty aesthetic: `width`.



Communicating PCA results

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[1]`) for the feature `concave.points_mean`?

```
# Loading of 'concave.points_mean' on PC1
wisc.pr$rotation["concave.points_mean", "PC1"]
```

```
[1] -0.2608538
```

The feature `concave.points_mean` has a loading of -0.2608538 on the first principal component (PC1), indicating it has low influence on PC1. The negative sign means that higher values of `concave.points_mean` tend to decrease the PC1 score for an observation. Since the loading of `concave.points_mean` on PC1 is -0.2608538 , higher values of this feature push PC1 toward the negative direction, which corresponds to malignant cases.

Q10. What is the minimum number of principal components required to explain 80% of the variance of the data?

```
# Cumulative proportion  
cum.pve <- cumsum(pve)  
  
# Find minimum number of PCs to reach 80%  
which(cum.pve >= 0.80) [1]
```

[1] 5

5 PCs are needed to explain 80% of the variance of the data.

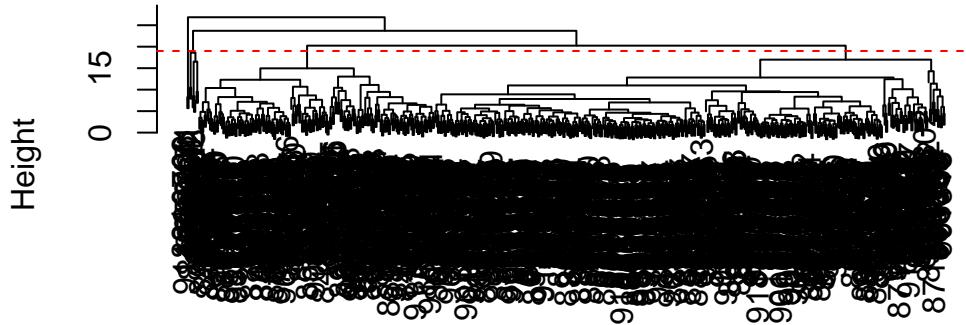
Hierarchical clustering

```
# Scale the wisc.data data using the "scale()" function  
data.scaled <- scale(wisc.data)  
data.dist <- dist(data.scaled)  
  
# Hierarchical clustering using complete linkage  
wisc.hclust <- hclust(data.dist, method = "complete")
```

Q11. Using the plot() and abline() functions, what is the height at which the clustering model has 4 clusters?

```
plot(wisc.hclust)  
abline(h=19, col="red", lty=2)
```

Cluster Dendrogram



```
data.dist  
hclust (*, "complete")
```

Height = 19 for when the clustering model has 4 clusters.

```
wisc.hclust.clusters <- cutree(wisc.hclust, k = 4)  
table(wisc.hclust.clusters, diagnosis)
```

wisc.hclust.clusters	B	M
1	12	165
2	2	5
3	343	40
4	0	2

Q12. Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 10?

```
best_k <- 0  
best_match <- 0  
  
for(k in 2:10){  
  clusters <- cutree(wisc.hclust, k = k)  
  tbl <- table(clusters, diagnosis)  
  match <- sum(apply(tbl, 1, max)) / length(diagnosis)}
```

```

if(match > best_match){
  best_match <- match
  best_k <- k
}
}

cat("Best number of clusters:", best_k, "\n")

```

Best number of clusters: 9

```
cat("Matching proportion:", best_match, "\n")
```

Matching proportion: 0.9103691

```
wisc.hclust.clusters <- cutree(wisc.hclust, k = 9)
table(wisc.hclust.clusters, diagnosis)
```

wisc.hclust.clusters	diagnosis	
	B	M
1	12	86
2	0	79
3	0	3
4	331	39
5	2	0
6	12	0
7	0	2
8	0	2
9	0	1

Q13. Which method gives your favorite results for the same data.dist dataset?
Explain your reasoning.

```
# Hierarchical clustering using ward.D2
wisc.hclust.ward <- hclust(data.dist, method = "ward.D2")

# Cut tree into 2 clusters (since we have benign vs malignant)
ward.clusters <- cutree(wisc.hclust.ward, k = 2)

# Compare to actual diagnosis
table(ward.clusters, diagnosis)
```

```

diagnosis
ward.clusters   B    M
      1  20 164
      2 337  48

# Optionally, compute proportion of correct matches
match_proportion <- sum(apply(table(ward.clusters, diagnosis), 1, max)) / length(diagnosis)
match_proportion

[1] 0.8804921

```

- The ward.D2 method minimizes the variance within clusters at each step, which tends to produce tight, spherical clusters that are more balanced in size. Complete linkage produces compact clusters but can overemphasize outliers, so ward.D2 is preferred.

K-means clustering

```

wisc.km <- kmeans(data.dist, centers = 2, nstart = 20)
table(wisc.km$cluster, diagnosis)

diagnosis
      B    M
      1 337  78
      2  20 134

```

Combining Methods (PCA and Clustering)

Clustering the original data was not very productive. The PCA results looked promising. Here we combine these methods by clustering from our PCA results. In other words, “clustering in PC space”... not using the original data but our PCA data (PC1, PC2, PCn...).

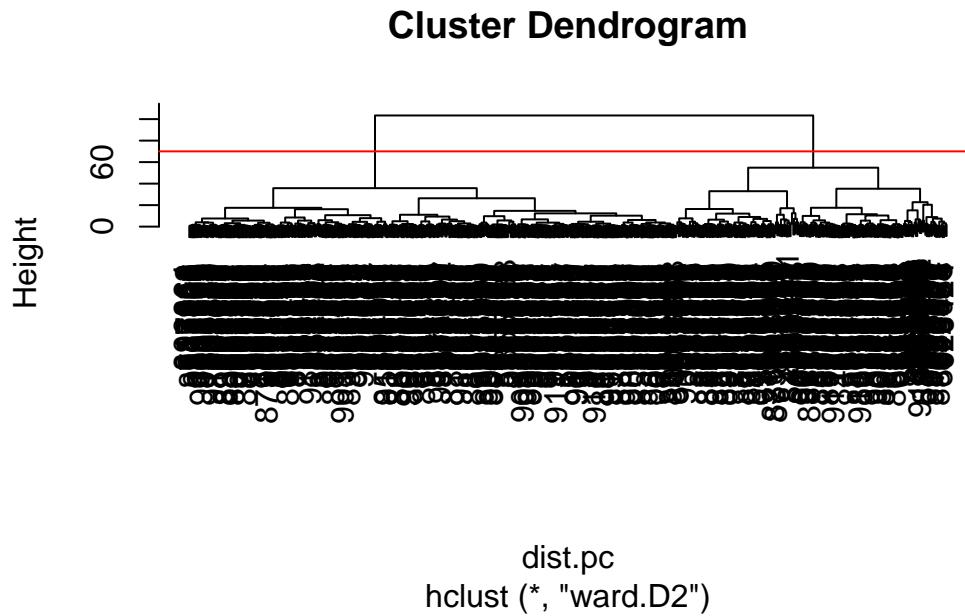
```

# Take the first 3 PCs
dist.pc <- dist(wisc.pr$x[,1:3])
wisc.pr.hclust <- hclust(dist.pc, method="ward.D2")

```

View the tree...

```
plot(wisc.pr.hclust)
abline(h=70, col="red")
```



To get our clustering membership vector (i.e. our main clustering result) we “cut” the tree at a desired height or to yield a desired number of “k” groups.

Q15. How well does the newly created model with four clusters separate out the two diagnoses? - Note for 15/16: we did not perform k-means clustering in class and this section was optional.

```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps, diagnosis)
```

diagnosis		
grps	B	M
1	24	179
2	333	33

Separates well. Separates into true positives = 179, false negatives = 33, false positives = 24, true negatives = 333.

Q16. How well do the k-means and hierarchical clustering models you created in previous sections (i.e. before PCA) do in terms of separating the diagnoses? Again, use the table() function to compare the output of each model (wisc.km\$cluster and wisc.hclust.clusters) with the vector containing the actual diagnoses. How does this clustering grps compare to the expert diagnosis?

```
table(wisc.km$cluster, diagnosis)
```

```
diagnosis
  B   M
1 337 78
2 20 134
```

We did not perform k-means cluster in class. However, on my own, k-means clustering separated into true positives = 134, false negatives = 78, false positives = 20, true negatives = 337. Separate similarly, although h-clustering + PCA does a better job with sensitivity, specificity is slightly larger for k-means clustering.

Q17. Which of your analysis procedures resulted in a clustering model with the best specificity? How about sensitivity?

We can calculate the sensitivity (how well h-clustering + PCA can detect positives/malignant).

Sensitivity (h-clustering + PCA): $TP/(TP+FN)$

```
179/(179+33)
```

```
[1] 0.8443396
```

We can calculate the specificity (how well h-clustering + PCA can detect negatives/benign).

Specificity (h-clustering + PCA): $TN/(TN+FP)$

```
333/(333 + 24)
```

```
[1] 0.9327731
```

Sensitivity (k-means clustering): $TP/(TP+FN)$

```
134/(134+78)
```

```
[1] 0.6320755
```

Specificity (k-means clustering): $TN/(TN+FP)$

```
337 / (337 + 20)
```

```
[1] 0.9439776
```

H-clustering + PCA has a higher sensitivity ($0.8443396 > 0.6320755$), while specificity is slightly larger for k-means clustering ($0.9439776 > 0.9327731$).

Prediction

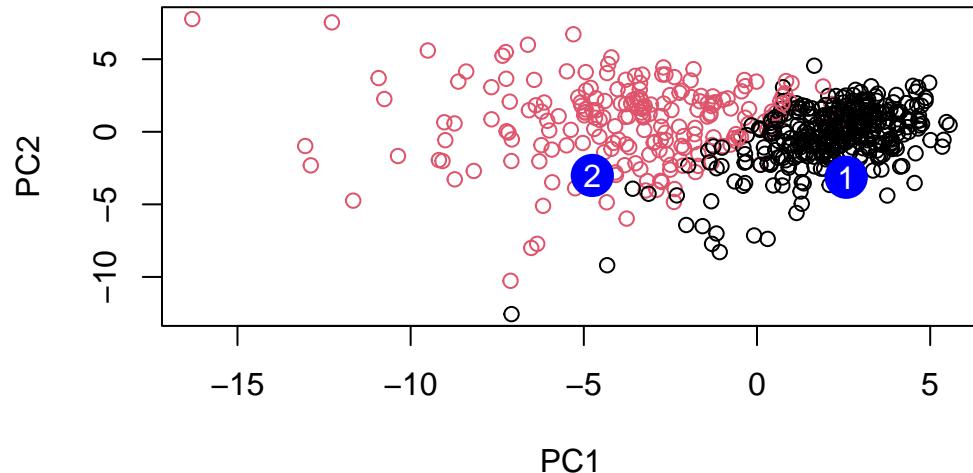
We can use our PCA model for prediction with new input patient samples.

```
#url <- "new_samples.csv"  
url <- "https://tinyurl.com/new-samples-CSV"  
new <- read.csv(url)  
npc <- predict(wisc.pr, newdata=new)  
npc
```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
[1,]	2.576616	-3.135913	1.3990492	-0.7631950	2.781648	-0.8150185	-0.3959098
[2,]	-4.754928	-3.009033	-0.1660946	-0.6052952	-1.140698	-1.2189945	0.8193031
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
[1,]	-0.2307350	0.1029569	-0.9272861	0.3411457	0.375921	0.1610764	1.187882
[2,]	-0.3307423	0.5281896	-0.4855301	0.7173233	-1.185917	0.5893856	0.303029
	PC15	PC16	PC17	PC18	PC19	PC20	
[1,]	0.3216974	-0.1743616	-0.07875393	-0.11207028	-0.08802955	-0.2495216	
[2,]	0.1299153	0.1448061	-0.40509706	0.06565549	0.25591230	-0.4289500	
	PC21	PC22	PC23	PC24	PC25	PC26	
[1,]	0.1228233	0.09358453	0.08347651	0.1223396	0.02124121	0.078884581	
[2,]	-0.1224776	0.01732146	0.06316631	-0.2338618	-0.20755948	-0.009833238	
	PC27	PC28	PC29	PC30			
[1,]	0.220199544	-0.02946023	-0.015620933	0.005269029			
[2,]	-0.001134152	0.09638361	0.002795349	-0.019015820			

```
g <- factor(diagnosis)
```

```
plot(wisc.pr$x[,1:2], col=g)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



Q18. Which of these new patients should we prioritize for follow up based on your results?

- Should follow up with patient 2.