

# Class 10: Halloween mini project

Joseph Girgiss (PID: A17388247)

## Table of contents

Data Import . . . . .	1
Quick overview of the dataset . . . . .	2
Winpercent and Pricepercent . . . . .	11
Exploring the correlation structure . . . . .	13
Principal component analysis . . . . .	14

As it is nearly Halloween and the half way point in the quarter let's do a mini project to help us figure out the best candy!

Our data comes from the 538 website and is available as a CSV file:

## Data Import

```
candy_file <- "candy-data.csv"  
candy = read.csv(candy_file, row.names=1)  
head(candy)
```

	chocolate	fruity	caramel	peanut	yalmond	nougat	crisped	rice	wafer
100 Grand	1	0	1		0	0			1
3 Musketeers	1	0	0		0	1			0
One dime	0	0	0		0	0			0
One quarter	0	0	0		0	0			0
Air Heads	0	1	0		0	0			0
Almond Joy	1	0	0		1	0			0
	hard	bar	pluribus	sugarpercent	pricepercent	winpercent			
100 Grand	0	1	0	0.732	0.860	66.97173			
3 Musketeers	0	1	0	0.604	0.511	67.60294			
One dime	0	0	0	0.011	0.116	32.26109			

```
One quarter      0   0       0       0.011      0.511    46.11650
Air Heads        0   0       0       0.906      0.511    52.34146
Almond Joy       0   1       0       0.465      0.767    50.34755
```

Q1. How many different candy types are in this dataset?

```
nrow(candy)
```

```
[1] 85
```

Q2. How many fruity candy types are in the dataset?

```
sum(candy$fruity)
```

```
[1] 38
```

```
candy["Twix", ]$winpercent
```

```
[1] 81.64291
```

Q3. What is your favorite candy in the dataset and what is it's winpercent value?

My favorite candy in the dataset is Milky Way.

```
candy["Milky Way", ]$winpercent
```

```
[1] 73.09956
```

Q4. What is the winpercent value for "Kit Kat"?

```
candy["Kit Kat", ]$winpercent
```

```
[1] 76.7686
```

Q5. What is the winpercent value for "Tootsie Roll Snack Bars"?

```
candy["Tootsie Roll Snack Bars", ]$winpercent
```

```
[1] 49.6535
```

## Quick overview of the dataset

```
library(skimr)
skimr::skim(candy)
```

Table 1: Data summary

Name	candy
Number of rows	85
Number of columns	12
Column type frequency:	
numeric	12
Group variables	None

### Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmondy	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

Yes, winpercent looks to be on a different scale since its values between 0-100 while all others are between 0 and 1.

Q7. What do you think a zero and one represent for the candy\$chocolate column?


```
candy$chocolate
```

```
[1] 1 1 0 0 0 1 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0 1 1 1 1 0 1 1 0 0 0 1 1 0 1 1  
[39] 1 1 1 0 1 1 0 0 0 1 0 0 0 1 1 1 1 0 1 0 0 1 0 0 1 0 1 1 0 0 0 0 0 0 0 1 1  
[77] 1 1 0 1 0 0 0 0 1
```

1 = the candy is a chocolate, 0 = the candy is not a chocolate.


Q8. Plot a histogram of winpercent values

```
library(ggplot2)  
  
ggplot(candy, aes(winpercent)) +  
  geom_histogram(binwidth = 5, fill = "steelblue", color = "white") +  
  labs(title = "Distribution of Candy Win Percentages",  
       x = "Win Percent", y = "Count") +  
  theme_bw()
```



Q9. Is the distribution of winpercent values symmetrical?

```
ggplot(candy, aes(winpercent)) +  
  geom_density() +  
  theme_bw()
```



The distribution of winpercent values is asymmetrical.

Q10. Is the center of the distribution above or below 50%?

```
mean(candy$winpercent)
```

[1] 50.31676

```
summary(candy$winpercent)
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	22.45	39.14	47.83	50.32	59.86	84.18

Center: mean = 50.31676 > 50%, median = 47.83 < 50% (indicating a right skew).

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

```
# 1. Find all chocolate candy in the dataset
# 2. Find their winpercent values
# 3. Calculate the mean of these values
# 4. Repeat for fruity candy.
# 5. Compare mean winpercents of chocolate vs fruity
```

```
choc inds <- candy$chocolate==1  
choc.win <- candy[choc inds, ]$winpercent  
choc.mean <- mean(choc.win)  
choc.mean
```

```
[1] 60.92153
```

```
fruit.inds <- candy$fruity==1  
fruit.win <- candy[fruit.inds, ]$winpercent  
fruit.mean <- mean(fruit.win)  
fruit.mean
```

```
[1] 44.11974
```

```
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
candy %>%  
  summarize(  
    mean_choc = mean(winpercent[chocolate == 1]),  
    mean_fruit = mean(winpercent[fruity == 1])  
)
```

```
mean_choc mean_fruit  
1 60.92153 44.11974
```

On average the chocolate candy is higher than fruity candy.

Q12. Is this difference statistically significant?

```
t.test(choc.win, fruit.win)
```

### Welch Two Sample t-test

```
data: choc.win and fruit.win
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 11.44563 22.15795
sample estimates:
mean of x mean of y
60.92153 44.11974
```

The difference between chocolate and fruity candy is statistically significant because the p-value is very small (p-value = 2.871e-08).

Q13. What are the five least liked candy types in this set?

```
candy %>%
  arrange(winpercent) %>%
  head(5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat	crispedrice	wafer	hard	bar	pluribus	sugarpercent	pricepercent
Nik L Nip	0	1	0	0	0	0	0	0	0	0	0	0.197	0.976
Boston Baked Beans	0	0	0	0	1	0	0	0	0	0	0	0.313	0.511
Chiclets	0	1	0	0	0	0	0	0	0	0	0	0.046	0.325
Super Bubble	0	1	0	0	0	0	0	0	0	0	0	0.162	0.116
Jawbusters	0	1	0	0	0	0	0	0	0	0	0	0.093	0.511
	winpercent												
Nik L Nip	22.44534												
Boston Baked Beans	23.41782												
Chiclets	24.52499												
Super Bubble	27.30386												
Jawbusters	28.12744												

Q14. What are the top 5 all time favorite candy types out of this set?

```
candy %>%
  arrange(desc(winpercent)) %>%
  head(5)
```


	chocolate	fruity	caramel	peanuty	almondy	nougat
Reese's Peanut Butter cup	1	0	0		1	0
Reese's Miniatures	1	0	0		1	0
Twix	1	0	1		0	0
Kit Kat	1	0	0		0	0
Snickers	1	0	1		1	1
	crispedrice	wafers	hard bar	pluribus	sugar	percent
Reese's Peanut Butter cup	0	0	0	0	0.720	
Reese's Miniatures	0	0	0	0	0.034	
Twix	1	0	1	0	0.546	
Kit Kat	1	0	1	0	0.313	
Snickers	0	0	1	0	0.546	
	price	percent	winpercent			
Reese's Peanut Butter cup	0.651	84.18029				
Reese's Miniatures	0.279	81.86626				
Twix	0.906	81.64291				
Kit Kat	0.511	76.76860				
Snickers	0.651	76.67378				

Q15. Make a first barplot of candy ranking based on winpercent values.

```
library(ggplot2)
candy$CandyName <- factor(rownames(candy), levels = rownames(candy))

ggplot(candy) +
  aes(x = winpercent, y = CandyName) +
  geom_col() +
  theme_bw() +
  labs(x = "Win Percent", y = "Candy", title = "Candy Rankings by Win Percent") +
  theme(
    axis.text.y = element_text(size = 4)
  )
```

## Candy Rankings by Win Percent




Q16. This is quite ugly, use the reorder() function to get the bars sorted by winpercent?

Let's also add some color based on the type of candy.

```
my_cols=rep("black", nrow(candy))
my_cols[as.logical(candy$chocolate)] = "chocolate"
my_cols[as.logical(candy$bar)] = "brown"
my_cols[as.logical(candy$fruity)] = "pink"

candy$CandyName <- factor(rownames(candy), levels = rownames(candy)[order(candy$winpercent)])
ggplot(candy) +
  aes(x = winpercent, y = CandyName) +
  geom_col(fill=my_cols) +
  theme_bw() +
  labs(x = "Win Percent", y = "Candy", title = "Candy Rankings by Win Percent") +
  theme(
    axis.text.y = element_text(size = 4)
  )
```

## Candy Rankings by Win Percent



Q17. What is the worst ranked chocolate candy?

```
candy %>%
  filter(chocolate == 1) %>%
  arrange(winpercent) %>%
  head(1)
```

	chocolate	fruity	caramel	peanut	almond	nougat	crisped rice	wafers	hard
Sixlets	1	0	0	0	0	0	0	0	0
	bar	pluribus	sugar	percent	price	percent	winpercent	CandyName	
Sixlets	0	1	0.22	0.081	34.722	0	34.722	Sixlets	0

Sixlets are the worst ranked chocolate candy.

Q18. What is the best ranked fruity candy?

```
candy %>%
  filter(fruity == 1) %>%
  arrange(desc(winpercent)) %>%
  head(1)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat	crisped	rice	wafer	hard
Starburst	0	1	0	0	0	0	0	0	0	0
	bar	pluribus	sugar	percent	price	percent	win	percent	CandyName	
Starburst	0	1	0.151		0.22	67.03763	Starburst			

Starburst is the best ranked fruity candy.


## Winpercent and Pricepercent

A plot with both variable/columns winpercent and pricepercent

```
library(ggrepel)

# How about a plot of price vs win
ggplot(candy) +
  aes(winpercent, pricepercent, label=rownames(candy)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols, size=3.3, max.overlaps = 5)
```

Warning: ggrepel: 65 unlabeled data points (too many overlaps). Consider increasing max.overlaps



Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

```
library(dplyr)
candy %>%
  mutate(Value = winpercent / (pricepercent + 1e-6)) %>%
  arrange(desc(Value)) %>%
  head(1)
```

```
      chocolate fruity caramel peanutyalmondynougat
Tootsie Roll Midgies      1       0       0           0       0
                           crispedricewafer hard bar pluribus sugarpercent
Tootsie Roll Midgies          0       0       0           1       0.174
                               pricepercent winpercent           CandyName   Value
Tootsie Roll Midgies     0.011    45.73675 Tootsie Roll Midgies 4157.508
```

Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

```
top5_expensive <- candy %>%
  arrange(desc(pricepercent)) %>% # sort by price descending
  slice(1:5) %>% # take top 5
  select(pricepercent, winpercent) # show relevant columns

top5_expensive
```

```
      pricepercent winpercent
Nik L Nip                  0.976  22.44534
Nestle Smarties             0.976  37.88719
Ring pop                   0.965  35.29076
Hershey's Krackel          0.918  62.28448
Hershey's Milk Chocolate    0.918  56.49050
```

```
least_popular <- top5_expensive %>%
  slice_min(winpercent) # candy with lowest winpercent

least_popular
```

```
      pricepercent winpercent
Nik L Nip                  0.976  22.44534
```


## Exploring the correlation structure

Now that we've explored the dataset a little, we'll see how the variables interact with one another. We'll use correlation and view the results with the corrplot package to plot a correlation matrix.

```
library(corrplot)
```

```
corrplot 0.95 loaded
```

```
cij <- cor(sapply(candy, is.numeric))
corrplot(cij)
```



Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?

```
# Convert the correlation matrix to a data frame for easier filtering
corr_df <- as.data.frame(as.table(cij))
colnames(corr_df) <- c("Variable1", "Variable2", "Correlation")

# Filter out the diagonal (where Variable1 equals Variable2)
off_diagonal_corr <- subset(corr_df, Variable1 != Variable2)
```

```
# Find the row with the minimum correlation value
most_anti_correlated <- off_diagonal_corr[which.min(off_diagonal_corr$Correlation), ]  
  
# Print the result
print("Most Anti-Correlated Pair:")
```

[1] "Most Anti-Correlated Pair:"

```
print(most_anti_correlated)
```

	Variable1	Variable2	Correlation
2	fruity	chocolate	-0.7417211

Fruity and Chocolate are most anti-correlated.

Q23. Similarly, what two variables are most positively correlated?

```
# Find the row with the maximum correlation value
most_positively_correlated <- off_diagonal_corr[which.max(off_diagonal_corr$Correlation), ]
```

```
# Print the result
print("Most Positively Correlated Pair:")
```

[1] "Most Positively Correlated Pair:"

```
print(most_positively_correlated)
```

	Variable1	Variable2	Correlation
12	winpercent	chocolate	0.6365167

Winpercent and chocolate are the most positively correlated pair.

## Principal component analysis

The function to use is called `prcomp()` with an optional `scale=T/F` argument.

```
candy_numeric <- candy[sapply(candy, is.numeric)]
pca <- prcomp(candy_numeric, scale=TRUE)
summary(pca)
```


Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369
	PC8	PC9	PC10	PC11	PC12		
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760		
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317		
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000		

Our main PCA result figure

```
ggplot(pca$x) +
  aes(PC1, PC2) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols, label=rownames(candy)) +
  theme_bw()
```

Warning: ggrepel: 48 unlabeled data points (too many overlaps). Consider increasing max.overlaps




We should also examine the variable “loadings” or contributions of the original variable to the new PCs.

```
pca$rotation
```


	PC1	PC2	PC3	PC4	PC5
chocolate	-0.4019466	0.21404160	0.01601358	-0.016673032	0.066035846
fruity	0.3683883	-0.18304666	-0.13765612	-0.004479829	0.143535325
caramel	-0.2299709	-0.40349894	-0.13294166	-0.024889542	-0.507301501
peanutyalmondy	-0.2407155	0.22446919	0.18272802	0.466784287	0.399930245
nougat	-0.2268102	-0.47016599	0.33970244	0.299581403	-0.188852418
crispedricewafer	-0.2215182	0.09719527	-0.36485542	-0.605594730	0.034652316
hard	0.2111587	-0.43262603	-0.20295368	-0.032249660	0.574557816
bar	-0.3947433	-0.22255618	0.10696092	-0.186914549	0.077794806
pluribus	0.2600041	0.36920922	-0.26813772	0.287246604	-0.392796479
sugarpcent	-0.1083088	-0.23647379	-0.65509692	0.433896248	0.007469103
pricepercent	-0.3207361	0.05883628	-0.33048843	0.063557149	0.043358887
winpercent	-0.3298035	0.21115347	-0.13531766	0.117930997	0.168755073
	PC6	PC7	PC8	PC9	PC10
chocolate	-0.09018950	-0.08360642	-0.49084856	-0.151651568	0.107661356
fruity	-0.04266105	0.46147889	0.39805802	-0.001248306	0.362062502
caramel	-0.40346502	-0.44274741	0.26963447	0.019186442	0.229799010
peanutyalmondy	-0.09416259	-0.25710489	0.45771445	0.381068550	-0.145912362

nougat	0.09012643	0.36663902	-0.18793955	0.385278987	0.011323453
crispedricewafer	-0.09007640	0.13077042	0.13567736	0.511634999	-0.264810144
hard	-0.12767365	-0.31933477	-0.38881683	0.258154433	0.220779142
bar	0.25307332	0.24192992	-0.02982691	0.091872886	-0.003232321
pluribus	0.03184932	0.04066352	-0.28652547	0.529954405	0.199303452
sugarpercent	0.02737834	0.14721840	-0.04114076	-0.217685759	-0.488103337
pricepercent	0.62908570	-0.14308215	0.16722078	-0.048991557	0.507716043
winpercent	-0.56947283	0.40260385	-0.02936405	-0.124440117	0.358431235
	PC11	PC12			
chocolate	0.10045278	0.69784924			
fruity	0.17494902	0.50624242			
caramel	0.13515820	0.07548984			
peanutyalmondy	0.11244275	0.12972756			
nougat	-0.38954473	0.09223698			
crispedricewafer	-0.22615618	0.11727369			
hard	0.01342330	-0.10430092			
bar	0.74956878	-0.22010569			
pluribus	0.27971527	-0.06169246			
sugarpercent	0.05373286	0.04733985			
pricepercent	-0.26396582	-0.06698291			
winpercent	-0.11251626	-0.37693153			

```
ggplot(pca$rotation) +
  aes(PC1, rownames(pca$rotation)) +
  geom_col()
```



```
plot(pca$x[,1:2], col=my_cols, pch=16)
```




```

# Make a new data-frame with our PCA results and candy data
my_data <- cbind(candy, pca$x[,1:3])

p <- ggplot(my_data) +
  aes(x=PC1, y=PC2,
      size=winpercent/100,
      text=rownames(my_data),
      label=rownames(my_data)) +
  geom_point(col=my_cols)

p

```



```

library(ggrepel)


p + geom_text_repel(size=3.3, col=my_cols, max.overlaps = 7) +
  theme(legend.position = "none") +
  labs(title="Halloween Candy PCA Space",
       subtitle="Colored by type: chocolate bar (dark brown), chocolate other (light brown),",
       caption="Data from 538")

```

Warning: ggrepel: 59 unlabeled data points (too many overlaps). Consider increasing max.overlaps

## Halloween Candy PCA Space

Colored by type: chocolate bar (dark brown), chocolate other (light brown),




Data from 538

Interactive plots that can be zoomed on and “brushed” over can be made with the **plotly** package. It’s output is interactive and will not render to PDF.

```
#library(plotly)
#ggplotly(p)

par(mar=c(8,4,2,2))
barplot(pca$rotation[,1], las=2, ylab="PC1 Contribution")
```



Q24. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you?

Fruity, hard, and pluribus are picked up strongly by PC1 in the positive direction. This makes sense because PC1 can be considered as a comparison of Chocolate vs fruity. It is logical that the features defining a typical non-chocolate/hard/many-pieces candy (Fruity, Pluribus, Hard) are strongly aligned in the positive direction of PC1, contrasting with the chocolate/bar/high-value (high pricepercent and winpercent) features in the negative direction.